

ARYAMAN MISHRA

19BCE1027

Lab Exercise 6

Take any text corpora, apply necessary preprocessing and perform the k-means clustering on the corpora.

Proposed Algorithm/Pseudocode:

Kmeans algorithm is an iterative algorithm that tries to partition the dataset into K pre-defined distinct non-overlapping subgroups (clusters) where each data point belongs to **only one group**. It tries to make the intra-cluster data points as similar as possible while also keeping the clusters as different (far) as possible. It assigns data points to a cluster such that the sum of the squared distance between the data points and the cluster's centroid (arithmetic mean of all the data points that belong to that cluster) is at the minimum. The less variation we have within clusters, the more homogeneous (similar) the data points are within the same cluster.

The way kmeans algorithm works is as follows:

1. Specify number of clusters K .
2. Initialize centroids by first shuffling the dataset and then randomly selecting K data points for the centroids without replacement.
3. Keep iterating until there is no change to the centroids. i.e assignment of data points to clusters isn't changing.
 - Compute the sum of the squared distance between data points and all centroids.
 - Assign each data point to the closest cluster (centroid).

- Compute the centroids for the clusters by taking the average of the all data points that belong to each cluster.

Data Structure Proposed: 2D-Arrays, Dictionaries.

IMPLEMENTATION CODE AND RESULTS:

```
import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
from sklearn.cluster import KMeans
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.decomposition import PCA
from sklearn.preprocessing import normalize
from sklearn.metrics import pairwise_distances
from nltk.tokenize import word_tokenize
from nltk.stem.porter import PorterStemmer
from nltk.corpus import stopwords
from bs4 import BeautifulSoup
from scipy.stats import multivariate_normal as mvn
import nltk
import os
import random
import string
import matplotlib.pyplot as plt
%matplotlib inline
plt.style.use('fivethirtyeight')

import os, sys, email, re
print(os.listdir("../input"))
```

```
['emails.csv']
```

```
df = pd.read_csv('../input/emails.csv', nrows = 35000)
df.shape
```

```
(35000, 2)
```

```
emails = list(map(email.parser.Parser().parsestr, df['message']))
headings = emails[0].keys()
for key in headings:
    df[key] = [doc[key] for doc in emails]
```

```

def get_raw_text(emails):
    email_text = []
    for email in emails.walk():
        if email.get_content_type() == 'text/plain':
            email_text.append(email.get_payload())
    return ''.join(email_text)


df['body'] = list(map(get_raw_text, emails))
df.head()
df['user'] = df['file'].map(lambda x: x.split('/')[0])

df['Date'] = pd.to_datetime(df['Date'], infer_datetime_format=True)
df.head()
df.dtypes

```

file	object
message	object
Message-ID	object
Date	datetime64[ns]
From	object
To	object
Subject	object
Mime-Version	object
Content-Type	object
Content-Transfer-Encoding	object
X-From	object
X-To	object
X-cc	object
X-bcc	object
X-Folder	object
X-Origin	object
X-FileName	object
body	object
user	object
dtype:	object


```
df['Month'] = df['Date'].dt.month
df['Year'] = df['Date'].dt.year
df['Day'] = df['Date'].dt.dayofweek
indices = (df['Year'] > 1995) & (df['Year'] <= 2004)
plt.figure(figsize = (10,6))
figure1 = df.loc[indices].groupby('Year')['body'].count().plot()
```




```
df.Year.agg({'max': max, 'min': min})
# really should not be dates up to 2044
df[df['Year']==2044]
```

```
file message Message-ID Date From To Subject Mime-Version Content-Type Content-Transfer-Encoding X-From X-To X-cc X-bcc X-Folder X-Origin X-FileName body user Month Year Day
```

```
plt.figure(figsize = (10,6))
figure2 = df.groupby('Month')['body'].count().plot()
```



```
plt.figure(figsize = (10,6))
figure3 = df.groupby('Day')['body'].count().plot()
```



```
#Unique to and From
print('Total number of emails: %d' %len(df))
print('-----')
print('Number of unique received: %d '%df['To'].nunique())
print('-----')
print('Number of unique Sent: %d '%df['From'].nunique())
```

Total number of emails: 35000

Number of unique received: 4926


Number of unique Sent: 2151

```
top_10_frequent = df.groupby('user')['file'].count().sort_values(ascending = False)[:30]
top_10_frequent
```

```
user
beck-s      11830
bass-e      7823
arnold-j     4898
allen-p     3034
baughman-d  2760
blair-l     1879
badeer-r     877
benson-r     767
arora-h      654
bailey-s     478
Name: file, dtype: int64
```

```
plt.figure(figsize = (10,6))
top_10_frequent.plot(kind = 'bar')
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f246e15f208>



```
df.groupby(['user', 'Year'])['file'].count()
```

user	Year	
allen-p	1980	10
	1999	8
	2000	1307
	2001	1704
arnold-j	2002	5
	1980	8
	2000	1832
	2001	3056
arora-h	2002	2
	1980	2
	2000	62
	2001	580
badeer-r	2002	10
	2000	851
	2001	2
	2002	24
bailey-s	2000	1
	2001	127
	2002	350
	1980	6
bass-e	1999	81
	2000	5727
	2001	1534
	2002	473
	2004	2
baughman-d	1980	6
	2000	176
	2001	2145
	2002	433
beck-s	1980	16
	1999	94
	2000	7235
	2001	4095
benson-r	2002	390
	2001	646
blair-l	2002	121
	2001	1873
	2002	6

Name: file, dtype: int64

```
def split_data(data):
    if data is not None:
        temp = data.split(',')
        if len(temp) == 1:
            return 'Direct'
        else:
            return 'Multiple'
    else:
        return 'Empty'
df['Direct_or_multi'] = df['To'].apply(split_data)
```

```
df.groupby('user')['Direct_or_multi'].value_counts().sort_values(ascending=False)[:15]
```

user	Direct_or_multi	
beck-s	Direct	7595
bass-e	Direct	5900
arnold-j	Direct	4354
beck-s	Multiple	3960
allen-p	Direct	2631
bass-e	Multiple	1738
baughman-d	Direct	1327
	Multiple	1311
blair-l	Empty	723
	Multiple	656
badeer-r	Direct	544
arora-h	Direct	505
blair-l	Direct	500
arnold-j	Multiple	408
allen-p	Multiple	340

Name: Direct_or_multi, dtype: int64

```
def clean_column(data):
    if data is not None:
        stopwords_list = stopwords.words('english')
        #exclusions = ['RE:', 'Re:', 're:']
        #exclusions = '|'.join(exclusions)
        data = data.lower()
        data = re.sub('re:', '', data)
        data = re.sub('-', '', data)
        data = re.sub('_', '', data)
        # Remove data between square brackets
        data = re.sub('\[[^\]]*\]', '', data)
        # removes punctuation
        data = re.sub(r'[\^\w\s]', '', data)
        data = re.sub(r'\n', ' ', data)
        data = re.sub(r'[0-9]+', '', data)
        # strip html
        p = re.compile(r'<.*?>')
        data = re.sub(r"\ve", " have ", data)
        data = re.sub(r"can't", "cannot ", data)
        data = re.sub(r"n't", " not ", data)
        data = re.sub(r"I'm", "I am", data)
        data = re.sub(r" m ", " am ", data)
        data = re.sub(r"\re", " are ", data)
        data = re.sub(r"\d", " would ", data)
        data = re.sub(r"\ll", " will ", data)
        data = re.sub('forwarded by phillip k allenhouect on pm', '', data)
        data = re.sub(r"httpitcappscorpenroncomsrrsauthemaillinkaspidpage", "", data)

        data = p.sub('', data)
        if 'forwarded by:' in data:
            data = data.split('subject')[1]
            data = data.strip()
            return data
        return 'No Subject'

df['Subject_new'] = df['Subject'].apply(clean_column)
df['body_new'] = df['body'].apply(clean_column)
```



```
0             here is our forecast
1  traveling to have a business meeting takes the...
2             test successful way to go
3  randy    can you send me a schedule of the sala...
4             lets shoot for tuesday at
Name: body_new, dtype: object
```

```
for i in to_add:
    stopwords.add(i)
```

```
print(wordcloud)
plt.figure(figsize=(9,8))
fig = plt.figure(1)
plt.imshow(wordcloud)
plt.axis('off')
plt.show()
```

[illegible]

```
def tokenize_and_stem(text):
    # first tokenize by sentence, then by word to ensure that punctuation is caught as it's own token
    tokens = [word for sent in nltk.sent_tokenize(text) for word in nltk.word_tokenize(sent)]
    filtered_tokens = []
    # filter out any tokens not containing letters (e.g., numeric tokens, raw punctuation)
    for token in tokens:
        if re.search('[a-zA-Z]', token):
            filtered_tokens.append(token)
    stems = [stemmer.stem(t) for t in filtered_tokens]
    return stems
```

```

from sklearn.feature_extraction.text import TfidfVectorizer
data = df['body_new']
# data.head()

tfidf_vectorizer = TfidfVectorizer(stop_words = stopwords, tokenizer = tokenize_and_stem,
                                   max_features = 5000)

%time tfidf = tfidf_vectorizer.fit_transform(data)
tfidf_norm = normalize(tfidf)
tfidf_array = tfidf_norm.toarray()
pd.DataFrame(tfidf_array, columns=tfidf_vectorizer.get_feature_names()).head()

/opt/conda/lib/python3.6/site-packages/sklearn/feature_extraction/text.py:301: UserWarning: Your stop_words may be inconsistent with your preprocessing. Tokenizing the stop words generated tokens ['aren', 'couldn', 'didn', 'doesn', 'don', 'fu', 'h
'stop_words.' % sorted(inconsistent))
CPU times: user 6.99 s, sys: 236 ms, total: 7.18 s
Wall time: 7.19 s

```

	aaron	abc	abs	absc	absstring	ablength	abel	abelhouectect	ability	able	absence	absolutely	ac	academy	accept	acceptable	acceptance	accepted	access	accessible	accessing	accommodate	accommodations	accomplish	accomplishments	according
0	0.0	0.0					0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
1	0.0	0.0					0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0					0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	0.0	0.0					0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0
4	0.0	0.0					0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0

```

# initial_centroids = np.random.permutation(tfidf_array.shape[0])[:3]
# initial_centroids
# centroids = tfidf_array[initial_centroids]
# centroids.shape
# dist_to_centroid = pairwise_distances(tfidf_array, centroids, metric = 'euclidean')
# cluster_labels = np.argmin(dist_to_centroid, axis = 1)

```

```
class Kmeans:
```

```

    def __init__(self, k, seed = None, max_iter = 200):
        self.k = k
        self.seed = seed
        if self.seed is not None:
            np.random.seed(self.seed)
        self.max_iter = max_iter

```

```
    def initialise_centroids(self, data):
```

```

        initial_centroids = np.random.permutation(data.shape[0])[:self.
k]
        self.centroids = data[initial_centroids]

        return self.centroids

```

```
    def assign_clusters(self, data):
```

```

        if data.ndim == 1:
            data = data.reshape(-1, 1)

        dist_to_centroid = pairwise_distances(data, self.centroids, me
tric = 'euclidean')
        self.cluster_labels = np.argmin(dist_to_centroid, axis = 1)

```

```

        return self.cluster_labels

def update_centroids(self, data):

    self.centroids = np.array([data[self.cluster_labels == i].mean(
axis = 0) for i in range(self.k)])

    return self.centroids

def convergence_calculation(self):
    pass

def predict(self, data):

    return self.assign_clusters(data)

def fit_kmeans(self, data):
    """
    This function contains the main loop to fit the algorithm
    Implements initialise centroids and update_centroids
    according to max_iter

    self.centroids = self.initialise_centroids(data)

    # Main kmeans loop
    for iter in range(self.max_iter):

        self.cluster_labels = self.assign_clusters(data)
        self.centroids = self.update_centroids(data)
        if iter % 100 == 0:
            print("Running Model Iteration %d " %iter)
    print("Model finished running")
    return self

```

```

from sklearn.datasets import make_blobs
data = make_blobs(n_samples=200, n_features=2, centers=4, cluster_std=1.6, random_state=50)
points = data[0]
plt.scatter(data[0][:,0], data[0][:,1], c=data[1], cmap='viridis')
plt.xlim(-15,15)
plt.ylim(-15,15)


```

```

X = data[0]
X[2]

```

```
array([-2.30970265,  5.8496944 ])
```



```

temp_k = Kmeans(4, 1, 600)
temp_fitted = temp_k.fit_kmeans(X)
new_data = np.array([[1.066, -8.66],
                     [1.87876, -6.516],
                     [-1.59728965,  8.45369045],
                     [1.87876, -6.516]])
temp_fitted.predict(new_data)

```

```

Running Model Iteration 0
Running Model Iteration 100
Running Model Iteration 200
Running Model Iteration 300
Running Model Iteration 400
Running Model Iteration 500
Model finished running
array([2, 2, 1, 2])

```

```

sklearn_pca = PCA(n_components = 2)
Y_sklearn = sklearn_pca.fit_transform(tf_idf_array)
test_e = Kmeans(3, 1, 600)
%time fitted = test_e.fit_kmeans(Y_sklearn)
predicted_values = test_e.predict(Y_sklearn)

plt.scatter(Y_sklearn[:, 0], Y_sklearn[:, 1], c=predicted_values, s=50, cmap='viridis')


centers = fitted.centroids
plt.scatter(centers[:, 0], centers[:, 1], c='black', s=300, alpha=0.6);

```

```

Running Model Iteration 0
Running Model Iteration 100
Running Model Iteration 200
Running Model Iteration 300
Running Model Iteration 400
Running Model Iteration 500
Model finished running
CPU times: user 2.41 s, sys: 4 ms, total: 2.42 s
Wall time: 2.41 s

```



```

from sklearn.cluster import KMeans
n_clusters = 3
sklearn_pca = PCA(n_components = 2)
Y_sklearn = sklearn_pca.fit_transform(tf_idf_array)
kmeans = KMeans(n_clusters= n_clusters, max_iter=600, algorithm = 'auto')
%time fitted = kmeans.fit(Y_sklearn)
prediction = kmeans.predict(Y_sklearn)

plt.scatter(Y_sklearn[:, 0], Y_sklearn[:, 1],c=prediction ,s=50, cmap='viridis')


centers2 = fitted.cluster_centers_
plt.scatter(centers2[:, 0], centers2[:, 1],c='black', s=300, alpha=0.6);

```

```

CPU times: user 192 ms, sys: 320 ms, total: 512 ms
Wall time: 132 ms

```




```
number_clusters = range(1, 7)

kmeans = [KMeans(n_clusters=i, max_iter = 600) for i in number_clusters]
kmeans

score = [kmeans[i].fit(Y_sklern).score(Y_sklern) for i in range(len(kmeans))]
score = [i*-1 for i in score]


plt.plot(number_clusters, score)
plt.xlabel('Number of Clusters')
plt.ylabel('Score')
plt.title('Elbow Method')
plt.show()
```



```
def get_top_features_cluster(tf_idf_array, prediction, n_feats):
    labels = np.unique(prediction)
    dfs = []
    for label in labels:
        id_temp = np.where(prediction==label) # indices for each cluster
        x_means = np.mean(tf_idf_array[id_temp], axis = 0) # returns average score across cluster
        sorted_means = np.argsort(x_means)[::-1][:n_feats] # indices with top 20 scores
        features = tf_idf_vectorizer.get_feature_names()
        best_features = [(features[i], x_means[i]) for i in sorted_means]
        df = pd.DataFrame(best_features, columns = ['features', 'score'])
        dfs.append(df)
    return dfs
dfs = get_top_features_cluster(tf_idf_array, prediction, 20)
```


```
import seaborn as sns
plt.figure(figsize=(8,6))
sns.barplot(x = 'score' , y = 'features', orient = 'h' , data = dfs[0][:15])
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f2462f19a20>




```
plt.figure(figsize=(8,6))
sns.barplot(x = 'score' , y = 'features', orient = 'h' , data = dfs[1][:15])
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f2462f49eb8>




```
plt.figure(figsize=(8,6))
sns.barplot(x = 'score' , y = 'features', orient = 'h' , data = dfs[2][:15])
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f2468074dd8>




```
for i, df in enumerate(dfs):
    df.to_csv('df_'+str(i)+'.csv')
```

```
def plot_features(dfs):
    fig = plt.figure(figsize=(14,12))
    x = np.arange(len(dfs[0]))
    for i, df in enumerate(dfs):
        ax = fig.add_subplot(1, len(dfs), i+1)
        ax.set_title("Cluster: " + str(i), fontsize = 14)
        ax.spines["top"].set_visible(False)
        ax.spines["right"].set_visible(False)
        ax.set_frame_on(False)
        ax.get_xaxis().tick_bottom()
        ax.get_yaxis().tick_left()
        ax.ticklabel_format(axis='x', style='sci', scilimits=(-2,2))
        ax.barh(x, df.score, align='center', color='#40826d')
        yticks = ax.set_yticklabels(df.features)
    plt.show();
plot_features(dfs)
```



CONCLUSION:

K-Means Clustering for email txt corpora has been successfully implemented and executed.